

```
In [22]: %matplotlib inline
import matplotlib
import matplotlib.pyplot as plt
import numpy as np
import seaborn as sns; sns.set()
import pandas as pd
```

```
In [23]: data1 = pd.read_csv("Wells.csv")
```


```
In [24]: data1.shape
```

```
Out[24]: (55945, 22)
```

```
In [25]: data2=data1[['Well','GR']]
```


```
In [26]: sns.boxplot(x=data1['Well'],y=data1['DT'])
```

```
Out[26]: <matplotlib.axes._subplots.AxesSubplot at 0x22a43c3e550>
```




```
In [27]: sns.boxplot(x=data1['Well'],y=data1['RHOB'])
```

```
Out[27]: <matplotlib.axes._subplots.AxesSubplot at 0x22a448be128>
```



```
In [28]: sns.boxplot(x=data2['Well'],y=data2['GR'])
```


```
Out[28]: <matplotlib.axes._subplots.AxesSubplot at 0x22a44a2d6a0>
```



```
In [29]: data3=data2[(data2['GR']<250)]
```

```
In [30]: sns.boxplot(x=data3['Well'],y=data3['GR'])
```


```
Out[30]: <matplotlib.axes._subplots.AxesSubplot at 0x22a44b9b898>
```



```
In [31]: d=data1[data1['Well']<13]
```

```
In [32]: sns.boxplot(x=d['Well'],y=d['ROP'])
```


```
Out[32]: <matplotlib.axes._subplots.AxesSubplot at 0x22a40334748>
```



```
In [33]: D1=d[(d['WOB']>0)&(d['WOB']<25)]
```

```
In [34]: sns.boxplot(x=D1['Well'],y=D1['WOB'])
```


```
Out[34]: <matplotlib.axes._subplots.AxesSubplot at 0x22a404d5048>
```



```
In [35]: D2=d[(d['Torque']>0)]
```

```
In [36]: sns.boxplot(x=D2['Well'],y=D2['Torque'])
```

```
Out[36]: <matplotlib.axes._subplots.AxesSubplot at 0x22a3e9f0fd0>
```



```
In [126]: d_s=d[['Well','Depth','GR','DT','RHOB']]
          d_scut=d_s[['GR','DT','RHOB']]
```

```
In [127]: d.shape
```

```
Out[127]: (51877, 22)
```

```
In [128]: from scipy import stats
          d_scut=d_scut[(np.abs(stats.zscore(d_scut)) < 3).all(axis=1)]
```

```
In [129]: d_s=d_s.merge(d_scut,on=['GR','DT','RHOB'])
```

```
In [130]: from sklearn.cluster import KMeans
          kmeans = KMeans(n_clusters=3, init='k-means++', n_init=10, max_iter=300)
          kmeans.fit(d_s[['GR','DT','RHOB']])
          print(kmeans.labels_)
          print(kmeans.cluster_centers_)
          y_kmeans = kmeans.predict(d_s[['GR','DT','RHOB']])
```

```
[1 1 1 ... 2 2 2]
[[ 44.0424164   84.51406858   2.40045972]
 [ 11.53640007   67.03649884   2.54585863]
 [116.01160529  92.08457433   2.48356251]]
```

```
In [131]: original_ds=d_s.copy()
```

```
In [132]: d_s['Facies']=y_kmeans.tolist()
```

```
In [133]: MWD1=d[['Well','Depth','ROP','WOB','Torque','SurfRPM','DownP','Mudflow','ECD']]
          MWD1_cut=MWD1[['ROP','WOB','Torque','SurfRPM','DownP','Mudflow','ECD']]
```

```
In [134]: from scipy import stats
          MWD1_cut=MWD1_cut[(np.abs(stats.zscore(MWD1_cut)) < 3).all(axis=1)]
```

```
In [135]: MWD1=MWD1.merge(MWD1_cut,on=['ROP','WOB','Torque','SurfRPM','DownP','Mudflow','ECD'])
```

```
In [136]: d_s=d_s.merge(MWD1,on=['Well','Depth'])
```

```
In [137]: d_s.to_csv("Wells - Copy.csv",index=False)
```

```
In [141]: d_s = pd.read_csv("Wells - Copy.csv")
          d_s.shape
```

```
Out[141]: (72834, 13)
```

```
In [138]: #Train_Set
          d_train=d_s[d_s['Well']<10]
```

```
In [139]: #Test_Set
d_test=d_s[d_s['Well']>=10]
```

```
In [140]: LogData=d_s[['ROP','WOB','Torque','SurfRPM','DownP','Mudflow','ECD']].values
```

```
In [141]: RockType=d_s['Facies'].values
```

```
In [117]: from sklearn.preprocessing import StandardScaler
X = LogData
y = RockType
X = StandardScaler().fit_transform(X)

from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.4)
```

```
In [142]: X_train=d_train[['ROP','WOB','Torque','SurfRPM','DownP','Mudflow','ECD']].values
X_test=d_test[['ROP','WOB','Torque','SurfRPM','DownP','Mudflow','ECD']].values
y_train=d_train['Facies'].values
y_test=d_test['Facies'].values
```

```
In [143]: from sklearn.preprocessing import StandardScaler
X_train = StandardScaler().fit_transform(X_train)
X_test = StandardScaler().fit_transform(X_test)
```

```
In [144]: # Import the SVM library
from sklearn.svm import NuSVC
from sklearn.svm import SVC
classifier = SVC(C = 100, kernel = 'rbf') # Define the SVM model parameters
classifier.fit(X_train, y_train) # Fit a classifier to the training data using
y-labels corresponding to properties stored in X.
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py:578: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n\_samples, ), for example using ravel().  
y = column\_or\_1d(y, warn=True)

```
Out[144]: SVC(C=100, cache_size=200, class_weight=None, coef0=0.0,
decision_function_shape='ovr', degree=3, gamma='auto', kernel='rbf',
max_iter=-1, probability=False, random_state=None, shrinking=True,
tol=0.001, verbose=False)
```

```
In [182]: y_pred = classifier.predict(X_test) # Predict the labels for the text X data

# If you do a train/test split, create a confusion matrix.
from sklearn.metrics import classification_report, confusion_matrix
print(confusion_matrix(y_test,y_pred))
print(classification_report(y_test,y_pred))
```

```
[[4177 2001   40]
 [ 359 3527   35]
 [ 907  148   18]]

              precision    recall  f1-score   support

     0       0.77       0.67       0.72       6218
     1       0.62       0.90       0.74       3921
     2       0.19       0.02       0.03       1073


 avg / total       0.66       0.69       0.66      11212
```

```
In [146]: print('generalization performance: ', classifier.score(X_test,y_test))
```

```
generalization performance:  0.6887263646093471
```


```
In [170]: plt.scatter(X_test[:,1],X_test[:,2], c=y_pred, cmap='rainbow')
plt.xlabel('WOB')
plt.ylabel('Torque')
plt.colorbar()
```

```
Out[170]: <matplotlib.colorbar.Colorbar at 0x218d39c50f0>
```



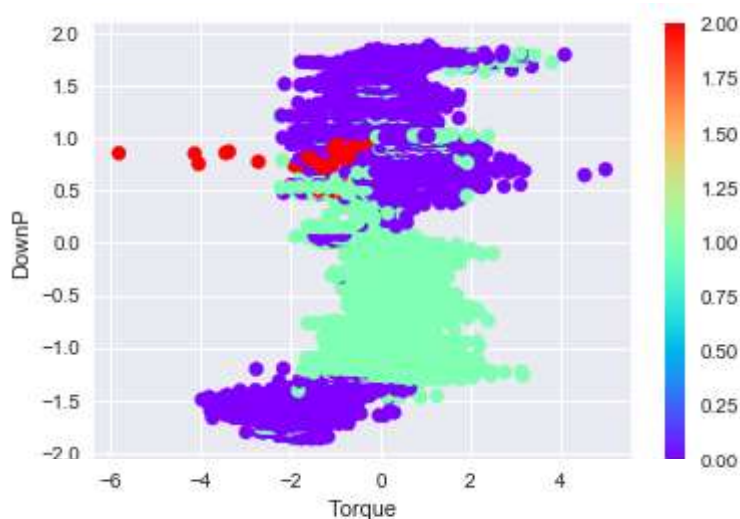
```
In [174]: plt.scatter(X_test[:,4],X_test[:,6], c=y_pred, cmap='rainbow')
plt.xlabel('DownP')
plt.ylabel('ECD')
plt.colorbar()
```

Out[174]: <matplotlib.colorbar.Colorbar at 0x218d4e12ef0>



```
In [179]: plt.scatter(X_test[:,2],X_test[:,4], c=y_pred, cmap='rainbow')
plt.xlabel('Torque')
plt.ylabel('DownP')
plt.colorbar()
```

Out[179]: <matplotlib.colorbar.Colorbar at 0x218d6f48cf8>



## RandomForest

```
In [231]: from sklearn.feature_selection import SelectFromModel
from sklearn.ensemble import RandomForestClassifier
r_clf=RandomForestClassifier(n_estimators = 200, max_depth=5)
```



```
In [232]: r_clf.fit(X_train,y_train)
```

C:\ProgramData\Anaconda3\lib\site-packages\ipykernel\_launcher.py:1: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n\_samples,), for example using ravel().

"""Entry point for launching an IPython kernel.

```
Out[232]: RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
                                max_depth=5, max_features='auto', max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, n_estimators=200, n_jobs=1,
                                oob_score=False, random_state=None, verbose=0,
                                warm_start=False)
```

```
In [233]: print('memorization performance: ', r_clf.score(X_train,y_train))

          print('generalization performance: ', r_clf.score(X_test,y_test))
```

```
memorization performance:  0.9106648792036282
generalization performance: 0.6073849447021049
```

```
In [234]: y_test_pred=r_clf.predict(X_test)
```

```
from sklearn.metrics import classification_report
print(classification_report(y_test, y_test_pred))
```

	precision	recall	f1-score	support
0	0.80	0.54	0.64	6218
1	0.49	0.88	0.63	3921
2	0.10	0.00	0.01	1073
avg / total	0.63	0.61	0.58	11212

```
In [235]: feature_list=list(d_train[['ROP','WOB','Torque','SurfRPM','DownP','Mudflow','ECD']].columns)
          feature_list
```

```
Out[235]: ['ROP', 'WOB', 'Torque', 'SurfRPM', 'DownP', 'Mudflow', 'ECD']
```

```
In [236]: # Get numerical feature importances
importances = list(r_clf.feature_importances_)
# List of tuples with variable and importance
feature_importances = [(feature, round(importance, 2)) for feature, importance
in zip(feature_list, importances)]
# Sort the feature importances by most important first
feature_importances = sorted(feature_importances, key = lambda x: x[1], reverse = True)
# Print out the feature and importances
[print('Variable: {:20} Importance: {}'.format(*pair)) for pair in feature_importances];
```

Variable: DownP	Importance: 0.32
Variable: ECD	Importance: 0.21
Variable: SurfRPM	Importance: 0.15
Variable: Torque	Importance: 0.13
Variable: Mudflow	Importance: 0.08
Variable: ROP	Importance: 0.06
Variable: WOB	Importance: 0.05


```
In [237]: import matplotlib.pyplot as plt
%matplotlib inline
# Set the style
plt.style.use('fivethirtyeight')
# list of x locations for plotting

x_values = list(range(len(importances)))

# Make a bar chart
plt.bar(x_values, importances, orientation = 'vertical')
# Tick labels for x axis

plt.xticks(x_values, feature_list, rotation='vertical')
# Axis labels and title

plt.ylabel('Importance'); plt.xlabel('Variable'); plt.title('Variable Importances');
```



```
In [238]: from sklearn.model_selection import KFold

cv = KFold(n_splits=3, shuffle = True, random_state=125)
```

```
In [239]: from sklearn.model_selection import GridSearchCV

param_grid = {'max_depth': [3,4,5,6,7,8], 'min_samples_split': [5,10,15]} # dictionary with keys

grid = GridSearchCV(r_clf, param_grid=param_grid, cv=cv, verbose=3)
```

```
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\model_selection\_validation.py:458: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().
```

```
    estimator.fit(X_train, y_train, **fit_params)
```

```
[CV] max_depth=8, min_samples_split=5, score=0.9366763648810542, total= 10.8s
```

```
[CV] max_depth=8, min_samples_split=10 .....
```

```
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\model_selection\_validation.py:458: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().
```

```
    estimator.fit(X_train, y_train, **fit_params)
```

```
[CV] max_depth=8, min_samples_split=10, score=0.9309489176007529, total= 9.2s
```

```
[CV] max_depth=8, min_samples_split=10 .....
```

```
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\model_selection\_validation.py:458: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().
```

```
    estimator.fit(X_train, y_train, **fit_params)
```

```
[CV] max_depth=8, min_samples_split=10, score=0.9358206400821496, total= 10.2s
```

```
[CV] max_depth=8, min_samples_split=10 .....
```

```
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\model_selection\_validation.py:458: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().
```

```
    estimator.fit(X_train, y_train, **fit_params)
```

```
[CV] max_depth=8, min_samples_split=10, score=0.9358206400821496, total= 9.1s
```

```
[CV] max_depth=8, min_samples_split=15 .....
```

```
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\model_selection\_validation.py:458: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().
```

```
    estimator.fit(X_train, y_train, **fit_params)
```

```
[CV] max_depth=8, min_samples_split=15, score=0.9314623085479593, total= 8.7s
```

```
[CV] max_depth=8, min_samples_split=15 .....
```

```
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\model_selection\_validation.py:458: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().
```

```
    estimator.fit(X_train, y_train, **fit_params)
```

```
[CV] max_depth=8, min_samples_split=15, score=0.936248502481602, total= 8.5s
```

```
[CV] max_depth=8, min_samples_split=15 .....
```

```
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\model_selection\_validation.py:458: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().
```

```
    estimator.fit(X_train, y_train, **fit_params)
```

```
[CV] max_depth=8, min_samples_split=15, score=0.9365052199212733, total=8.4s
```

```
[Parallel(n_jobs=1)]: Done 54 out of 54 | elapsed: 6.9min finished
```

```
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\model_selection\_search.py:739: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().
```

```
    self.best_estimator_.fit(X, y, **fit_params)
```

```
Out[240]: GridSearchCV(cv=KFold(n_splits=3, random_state=125, shuffle=True),
    error_score='raise',
    estimator=RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
    max_depth=5, max_features='auto', max_leaf_nodes=None,
    min_impurity_decrease=0.0, min_impurity_split=None,
    min_samples_leaf=1, min_samples_split=2,
    min_weight_fraction_leaf=0.0, n_estimators=200, n_jobs=1,
    oob_score=False, random_state=None, verbose=0,
    warm_start=False),
    fit_params=None, iid=True, n_jobs=1,
    param_grid={'max_depth': [3, 4, 5, 6, 7, 8], 'min_samples_split': [5,
    10, 15]},
    pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',
    scoring=None, verbose=3)
```

```
In [241]: print(grid.best_params_)
print("score", grid.score(X_test,y_test))

{'max_depth': 8, 'min_samples_split': 5}
score 0.6196039957188726
```

```
In [242]: print(grid.cv_results_.keys()) #for each parameter combination following metrics are stored

dict_keys(['mean_fit_time', 'std_fit_time', 'mean_score_time', 'std_score_time', 'param_max_depth', 'param_min_samples_split', 'params', 'split0_test_score', 'split1_test_score', 'split2_test_score', 'mean_test_score', 'std_test_score', 'rank_test_score', 'split0_train_score', 'split1_train_score', 'split2_train_score', 'mean_train_score', 'std_train_score'])
```

```
In [243]: import pandas as pd

cv_results = pd.DataFrame(grid.cv_results_)

# generate a subset of the table
cv_results_tiny = cv_results[['param_max_depth', 'param_min_samples_split', 'mean_test_score', 'std_test_score']]

# rank them based on test scores
cv_results_tiny.sort_values(by='mean_test_score', ascending=False).head(5)
```

```
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\deprecation.py:122:
FutureWarning: You are accessing a training score ('split0_train_score'), which will not be available by default any more in 0.21. If you need training scores, please set return_train_score=True
  warnings.warn(*warn_args, **warn_kwargs)
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\deprecation.py:122:
FutureWarning: You are accessing a training score ('split1_train_score'), which will not be available by default any more in 0.21. If you need training scores, please set return_train_score=True
  warnings.warn(*warn_args, **warn_kwargs)
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\deprecation.py:122:
FutureWarning: You are accessing a training score ('split2_train_score'), which will not be available by default any more in 0.21. If you need training scores, please set return_train_score=True
  warnings.warn(*warn_args, **warn_kwargs)
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\deprecation.py:122:
FutureWarning: You are accessing a training score ('mean_train_score'), which will not be available by default any more in 0.21. If you need training scores, please set return_train_score=True
  warnings.warn(*warn_args, **warn_kwargs)
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\deprecation.py:122:
FutureWarning: You are accessing a training score ('std_train_score'), which will not be available by default any more in 0.21. If you need training scores, please set return_train_score=True
  warnings.warn(*warn_args, **warn_kwargs)
```

Out[243]:

	param_max_depth	param_min_samples_split	mean_test_score	std_test_score
15	8	5	0.935338	0.002915
17	8	15	0.934739	0.002319
16	8	10	0.934197	0.002297
12	7	5	0.928150	0.002570
13	7	10	0.927950	0.002480

```
In [244]: grid_new=grid.best_estimator_
```

```
In [245]: print('memorization performance: ', grid_new.score(X_train,y_train)) # memorization
          print('generalization performance: ', grid_new.score(X_test,y_test)) # generalization

memorization performance:  0.9419264668130866
generalization performance:  0.6196039957188726
```

## PCA

```
In [326]: d_pca=d_s[['GR','DT','RHOB']]
```

```
In [327]: dpca=d_pca
```

```
In [328]: from sklearn.cluster import KMeans
          kmeans = KMeans(n_clusters=3, init='k-means++', n_init=10, max_iter=300)
          kmeans.fit(dpca)
          print(kmeans.labels_)
          print(kmeans.cluster_centers_)
          y_kmeans = kmeans.predict(dpca)
```

```
[1 1 1 ... 2 2 2]
[[ 43.67232042  83.97561779   2.40660651]
 [ 11.32876413  66.91932523   2.54839223]
 [116.71018132  91.53289838   2.48708824]]
```

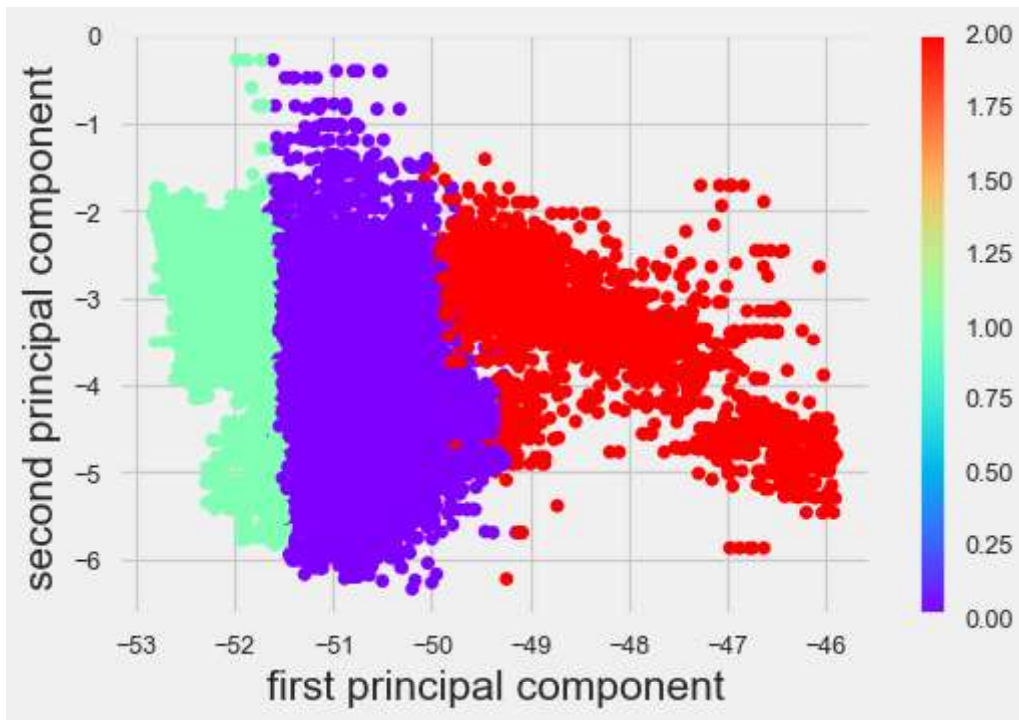
```
In [329]: from sklearn.decomposition import PCA
          pca = PCA(n_components=3)
```

```
In [330]: pca.fit(dpca)
```

```
Out[330]: PCA(copy=True, iterated_power='auto', n_components=3, random_state=None,
            svd_solver='auto', tol=0.0, whiten=False)
```

```
In [331]: from sklearn.preprocessing import StandardScaler
X_dpca=np.array(dpca)
X_dpca = StandardScaler().fit_transform(X_dpca)
X_pca = pca.transform(X_dpca)
plt.rcParams['figure.dpi']=90
plt.scatter(X_pca[:, 0], X_pca[:, 2], linewidths=0, s=30,c=y_kmeans,cmap='rain
bow')
plt.xlabel("first principal component")
plt.ylabel("second principal component")
plt.colorbar()
```

Out[331]: <matplotlib.colorbar.Colorbar at 0x218805095c0>





```

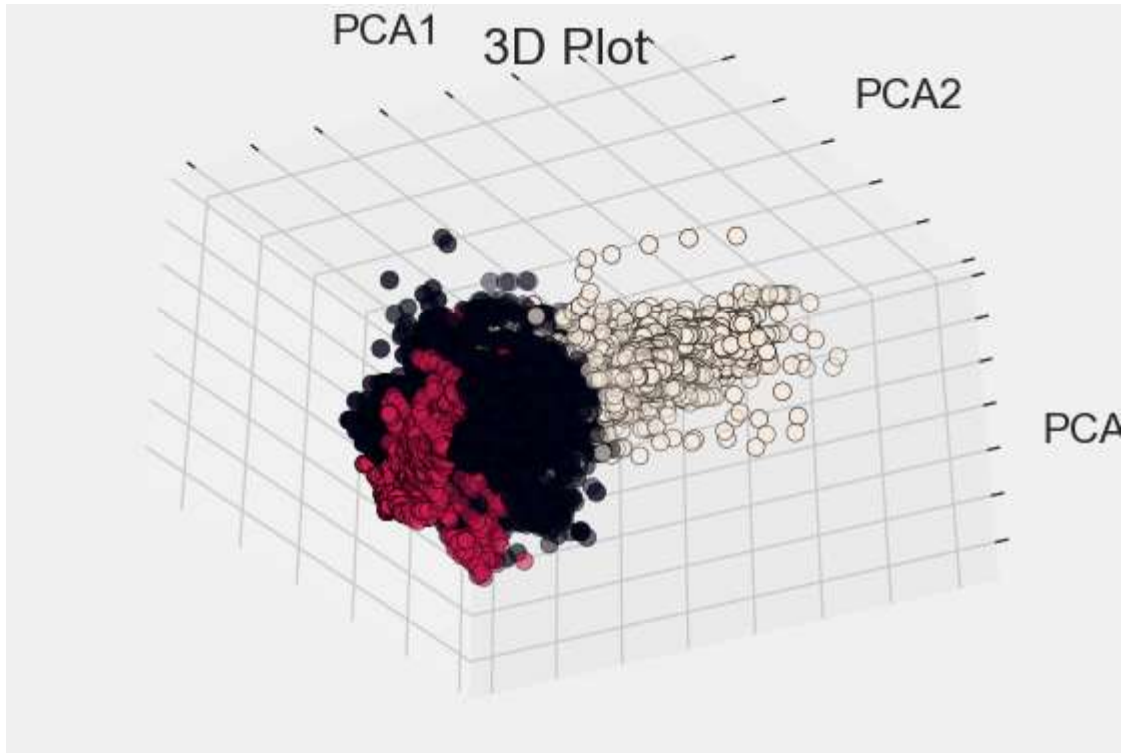
In [332]: from matplotlib.lines import Line2D
import matplotlib.pyplot as plt
from mpl_toolkits.mplot3d import Axes3D
import numpy as np
from sklearn import decomposition
from sklearn import datasets

fig = plt.figure(1)
ax = Axes3D(fig, elev=-40, azimuth=300) #3D plot
ax.scatter(X_pca[:, 0], X_pca[:, 1], X_pca[:, 2], c=y_kmeans, s = 50, edgecolor='k')

ax.w_xaxis.set_ticklabels([])
ax.w_yaxis.set_ticklabels([])
ax.w_zaxis.set_ticklabels([])
ax.set_xlabel('PCA1')
ax.set_ylabel('PCA2')
ax.set_zlabel('PCA3')
ax.set_title('3D Plot')

```

Out[332]: Text(0.5,0.92,'3D Plot')



```
In [333]: fig = plt.figure(figsize=(5.5, 3))
ax = Axes3D(fig, rect=[0, 0, .7, 1], elev=48, azimuth=134)
labelTups = [('PCA1', 0), ('PCA2', 1), ('PCA3', 2)]
for name, label in labelTups:
    ax.text3D(X[y == label, 0].mean(),
              X[y == label, 1].mean() + 1.5,
              X[y == label, 2].mean(), name,
              horizontalalignment='center',
              bbox=dict(alpha=.5, edgecolor='w', facecolor='w'))
# Reorder the labels to have colors matching the cluster results

sc = ax.scatter(X_pca[:, 0], X_pca[:, 1], X_pca[:, 2], c=y_kmeans, cmap="Spectral", edgecolor='k')

ax.w_xaxis.set_ticklabels([])
ax.w_yaxis.set_ticklabels([])
ax.w_zaxis.set_ticklabels([])

colors = [sc.cmap(sc.norm(i)) for i in [1, 2, 0]]
custom_lines = [plt.Line2D([], [], ls="", marker='.',
                           mec='k', mfc=c, mew=.1, ms=20) for c in colors]
ax.legend(custom_lines, [lt[0] for lt in labelTups],
          loc='center left', bbox_to_anchor=(1.0, .5))

plt.show()
```

